

Computer Vision HW2 report

Author: 109550093 黃得誠



I. Implementation

In this homework, I learnt how to stitch images using SIFT features and homography matrix with RANSAC. I implemented the algorithms using Python and OpenCV. I started from stitching 2 images. I used the right image as base, project left image onto right image coordinates. For baseline image cases, I stitch the rightmost first, from right to left, 2 images at a time. For the bonus images, I stitched m3 and m4, m1 and m2 first, then stitch the two stitched image together to form the final image.

1. SIFT

For SIFT, I simply followed the spec, using OpenCV functions to extract SIFT features. “kp” is OpenCV keypoints, which stores the point information.

“descript” is an array, which stores the gradient information.

```
# Step 1 SIFT
print("Step 1 SIFT ...")
kp_left, des_left = SIFT(img_left_g)
kp_right, des_right = SIFT(img_right_g)

24 # SIFT
25 def SIFT(img):
26     SIFT_Detector = cv2.SIFT_create()
27     kp, descript = SIFT_Detector.detectAndCompute(img, None)
28     return kp, descript
```

2. Feature Matching-kNN and Lowe's Ratio test

In step 2, I apply kNN algorithm (in a brutal Force way) and Lowe's Ratio test to find the good matches. I used OpenCV.BFMatcher() to debug. Since I found that OpenCV uses a class called dmatch to store the match information, I created dictionaries to store the information in a similar way.

```
440 # Step 2 Feature matching
441 print("Step 2 Feature Matching ...")
442 good, good1 = matcher(kp_left, des_left, kp_right, des_right, 0.7)
443 matches = get_matches(good1, kp_left, kp_right)
```

```
36 def matcher(kp1, des1, kp2, des2, threshold):
```

To find the least and second least distance, I created two dictionaries, dmatch and dmatch1 to store. I used np.linalg.norm to calculate the distance faster.

```
50 # Brutal Force
51 matches1 = []
52 for i in tqdm(range(len(kp1))):
53     dmatch = {"distance":1e7, "queryIdx":0, "trainIdx":0} # smallest
54     dmatch1 = {"distance":1e7, "queryIdx":0, "trainIdx":0} # 2nd smallest
55     v1 = des1[i,:]
56     for j in range(len(kp2)):
57         v2 = des2[j,:]
58         distance = 0
59         distance = np.linalg.norm(v1-v2)
60         if distance < dmatch["distance"]:
61             dmatch1["distance"] = dmatch["distance"]
62             dmatch1["queryIdx"] = dmatch["queryIdx"]
63             dmatch1["trainIdx"] = dmatch["queryIdx"]
64             dmatch["distance"] = distance
65             dmatch["queryIdx"] = i
66             dmatch["trainIdx"] = j
67         elif distance < dmatch1["distance"]:
68             dmatch1["distance"] = distance
69             dmatch1["queryIdx"] = i
70             dmatch1["trainIdx"] = j
71     matches1.append((dmatch, dmatch1))
```

Finally, the ratio test.

```
72
73 good1 = []
74 # Apply Lowe's ratio test
75 for m,n in matches1:
76     # m, n are DMatch
77     if m["distance"] < threshold*n["distance"]:
78         good1.append([m])
79
80 return good, good1
```

I stored the matches information in np.array "matches" for later use.

```
83 def get_matches(good, kp1, kp2):
84     ...
85     out: matches [[x1, y1, x1', y1'], ....., [xn, yn, xn', yn']]
86     ...
87     matches = []
88     for pair in good:
89         matches.append(list(kp1[pair[0]["queryIdx"]].pt + kp2[pair[0]["trainIdx"]].pt))
90         # matches.append(list(kp1[pair[0].queryIdx].pt + kp2[pair[0].trainIdx].pt))
91     matches = np.array(matches)
92     return matches
```

3. Homography matrix with RANSAC

For RANSAC, I followed the pseudocode in the spec. I decided the iterations and the error threshold by testing several times. I didn't apply a stopping criterion.

```
452 # Step 3 Homography
453 print("Step 3 RANSAC ...")
454 inliers, H = ransac(matches, 0.4, 8000)
```

```
149 def ransac(matches, threshold, iters):
150     num_best_inliers = 0
151     best_H = np.zeros((3, 3))
152     for i in tqdm(range(iters)):
153         points = random_point(matches)
154         H = homography(points)
155
156         # avoid dividing by zero
157         if np.linalg.matrix_rank(H) < 3:
158             continue
159
160         errors = get_error(matches, H)
161         idx = np.where(errors < threshold)[0]
162         inliers = matches[idx]
163
164         num_inliers = len(inliers)
165         if num_inliers > num_best_inliers:
166             best_inliers = inliers.copy()
167             num_best_inliers = num_inliers
168             best_H = H.copy()
169
170     print("inliers/matches: {}/{}".format(num_best_inliers, len(matches)))
171     return best_inliers, best_H
```

random_point and get_error are some helper functions that I created.

```
127 # Helper functions of RANSAC
128 def random_point(matches, k=4):
129     # Get k pairs from the matches (Used in RANSAC)
130     idx = random.sample(range(len(matches)), k)
131     point = [matches[i] for i in idx]
132     return np.array(point)
133
134 def get_error(matches, H):
135     # matches: [[x1, y1, x1', y1'], ..., [xn, yn, xn', yn']]
136     num_points = len(matches)
137     all_p1 = np.concatenate((matches[:, 0:2], np.ones((num_points, 1))), axis=1) # change (x1, y1) to (x1, y1, 1)
138     all_p2 = matches[:, 2:4]
139     estimate_p2 = np.zeros((num_points, 2))
140     for i in range(num_points):
141         temp = np.dot(H, all_p1[i])
142         estimate_p2[i] = (temp/temp[2])[0:2] # set index 2 to 1 and slice the index 0, 1
143     # Compute error
144     errors = np.linalg.norm(all_p2 - estimate_p2, axis=1) ** 2
145
146     return errors
```

Finally, the homography matrix.

```

110 # Step 3 Homography
111 def homography(pairs):
112     # pairs: [(x1, y1), (x1', y1')], [(x2, y2), (x2', y2')], [(x3, y3), (x3', y3')], [(x4, y4), (x4', y4')]
113     rows = []
114     for i in range(pairs.shape[0]):
115         p1 = np.append(pairs[i][0:2], 1)
116         p2 = np.append(pairs[i][2:4], 1)
117         row1 = [0, 0, 0, p1[0], p1[1], p1[2], -p2[1]*p1[0], -p2[1]*p1[1], -p2[1]*p1[2]]
118         row2 = [p1[0], p1[1], p1[2], 0, 0, 0, -p2[0]*p1[0], -p2[0]*p1[1], -p2[0]*p1[2]]
119         rows.append(row1)
120         rows.append(row2)
121     rows = np.array(rows) # A in spec, rows.shape is (8, 9)
122     U, s, V = np.linalg.svd(rows)
123     H = V[-1].reshape(3, 3) # use the last vector since np.linalg.svd has vector sorted in descending order
124     H = H/H[2, 2] # normalize h33 to 1
125     return H

```

I followed this page (with a minus sign) in the spec to calculate my homography matrix.

3. Homography

$$A = U\Sigma V^T$$

- Using **SVD decomposition** to find Least Squares error solution of **Ah = 0**
- the solution = eigenvector of $A^T A$ associated with the smallest eigenvalue (V stores the eigenvector of $A^T A$, Σ stores the singular value (root of eigen value))
find the **smallest number** in Σ and **H = corresponding vector in V^T**
- Remember to **normalize h33 to 1**

$$\begin{aligned}
 \mathbf{h} &= (H_{11}, H_{12}, H_{13}, H_{21}, H_{22}, H_{23}, H_{31}, H_{32}, H_{33})^T \\
 \mathbf{a}_x &= (-x_1, -y_1, -1, 0, 0, 0, x'_2 x_1, x'_2 y_1, x'_2)^T \\
 \mathbf{a}_y &= (0, 0, 0, -x_1, -y_1, -1, y'_2 x_1, y'_2 y_1, y'_2)^T
 \end{aligned}
 \quad A = \begin{pmatrix} \mathbf{a}_{x1}^T \\ \mathbf{a}_{y1}^T \\ \vdots \\ \mathbf{a}_{xN}^T \\ \mathbf{a}_{yN}^T \end{pmatrix}$$

You can multiply a minus to match the form in previous slide **A** is a 9 by 9 matrix (It's similar to A in previous slide)

Reference :

By the way, I think that A is 8x9 (since there should be even rows).

I used np.linalg.svd to solve the Least squares problem. Since the function give us eigenvalues in a descending order, I get my H through the last row.

Returns: **u** : { (... , M, M), (... , M, K) } array
Unitary array(s). The first **a.ndim - 2** dimensions have the same size as those of the input **a**. The size of the last two dimensions depends on the value of **full_matrices**. Only returned when **compute_uv** is True.

s : (... , K) array
Vector(s) with the singular values, within each vector sorted in descending order. The first **a.ndim - 2** dimensions have the same size as those of the input **a**.

vh : { (... , N, N), (... , K, N) } array
Unitary array(s). The first **a.ndim - 2** dimensions have the same size as those of the input **a**. The size of the last two dimensions depends on the value of **full_matrices**. Only returned when **compute_uv** is True.

4. Image Stitching

Last but not least, the image stitching part. I followed the spec as well.

```
457     # Step 4 Stitch Image
458     print("Step 4 stitching image ...")
459     img_stitched = stitch_img(img_left, img_right, H)
```

First, convert the image to double for further multiplication.

```
343     def stitch_img(left, right, H):
344
345         # Convert to double and normalize. Avoid noise.
346         left = cv2.normalize(left.astype('float'), None,
347                               0.0, 1.0, cv2.NORM_MINMAX)
348         # Convert to double and normalize.
349         right = cv2.normalize(right.astype('float'), None,
350                                0.0, 1.0, cv2.NORM_MINMAX)
```

Second, calculate the new corners. Calculate the new size and the affine matrix mentioned in the spec. Then do matrix multiplications by `cv2.warpPerspective`.

```
352     # left image
353     height_l, width_l, channel_l = left.shape
354     corners = [[0, 0, 1], [width_l, 0, 1], [width_l, height_l, 1], [0, height_l, 1]]
355     corners_new = [np.dot(H, corner) for corner in corners]
356     corners_new = np.array(corners_new).T
357     x_news = corners_new[0] / corners_new[2]
358     y_news = corners_new[1] / corners_new[2]
359     y_min = min(y_news)
360     x_min = min(x_news)
361     y_min = min(y_min, 0)
362     x_min = min(x_min, 0)
363
364     translation_mat = np.array([[1, 0, -x_min], [0, 1, -y_min], [0, 0, 1]])
```

Finally, concatenate the two images.



baseline/m4 and baseline/m5

For concatenating, I used blender method. I applied “linear blending” on the baseline images. Linear blending means that in the overlap area, I give each pixel different weights, if the pixel is closer to the left picture, the left picture weight will be higher. I eliminated some strange boundaries after applying blender method.



With and without blender.

I also wrote a removeBlackBorder function to eliminate the black border by going through the pixels.

By the way, in the baseline images, to make the result looks better, I cut the top of the image to avoid this kind of result happens.



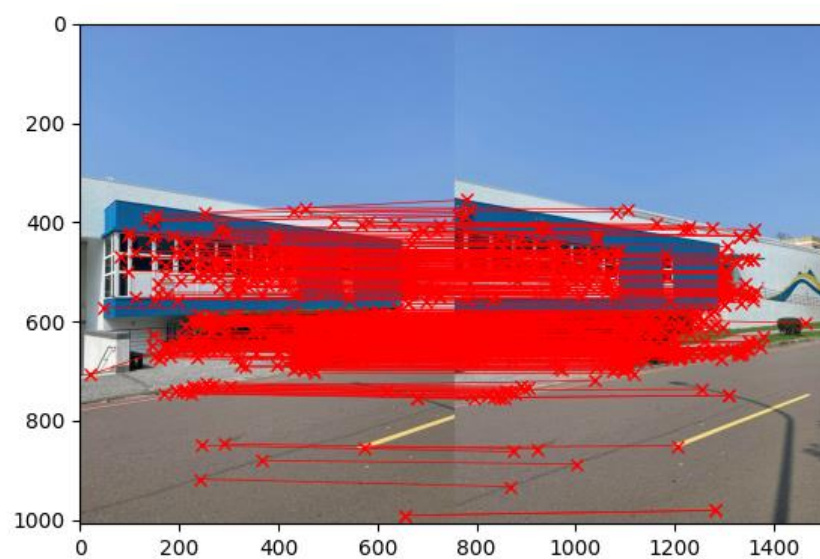
However, I didn't do it in the bonus, since the way that bonus images have taken is not in the same order. I change the stitching order instead in the bonus, as I mentioned earlier.

This is the end of the introduction of my implementation. Here are some figures that could show the results of each step.


SIFT:




Feature matching



RANSAC



Mask that blending method used.



II. Result of 2 images

1. Baseline



2. Bonus



III. Stitching more images

1. Baseline



2. Bonus



IV. Conclusion

In this homework, I implemented an image stitching method that have been taught in class. Although taught in class, I still met some problems when implementing.

The first problem is the boundary problem.



The alias problem and inverse mapping came to my mind when I met this problem, but I didn't use the inverse mapping to cope with it. Because after blending, the problem seems to be solved.

The second problem occurred when I'm stitching the bonus images.




There are still some boundaries that can be clearly seen after linear blending. I tried “linear blending with constant width” instead, which only blend constant width instead of all the overlapped parts.



It kind of solve some part of the problem but new boundaries occurred too. I think I need to spend more time to test with the constant to have better results, which I didn't continued finally. Since stitching 4 pictures could somehow cover those boundaries.



Last thing that I discovered is that linear blending with constant width can solve another problem, called ghost.



The ghost problem means that there are some scene that are doubled, due to the blending process. Linear blending with constant width doesn't blend that much so the ghost problem would not occur.

Finally, I learnt a lot in this homework.